

THE STORAGE CRASHCOURSE

Benjamin Ritter



- Platform Engineering Consultant @ Scandio
- Working mostly with Cloud Native Technologies



l0wl3vel





You will see the asterisk a lot

I will generalize and simplify some concepts. Things will be technically wrong, but still useful to understand the rough concept

K8S STORAGE RESOURCE MODEL

Kubernetes does not contain a baked in solution for storage, only interfaces to integrate with external storage providers

instead it allows you to bring your own via the
Container Storage Interface

PERSISTENT VOLUME CLAIM

- Namespaced
- Contains a volume “request”
- Attached to pod
- The resource you will be working with most

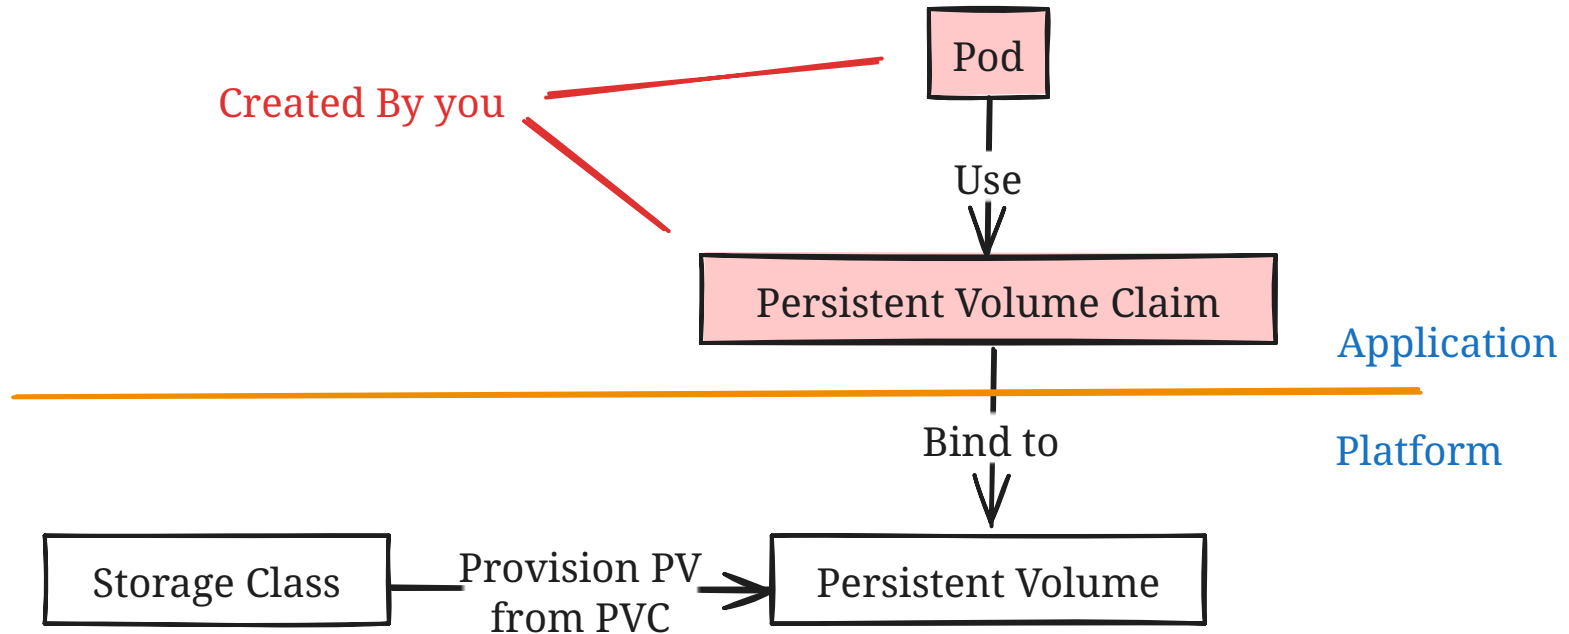
PERSISTENT VOLUME

- Reference to the real volume on the storage platform
- Contains implementation details for the used storage driver

STORAGE CLASS (SC):

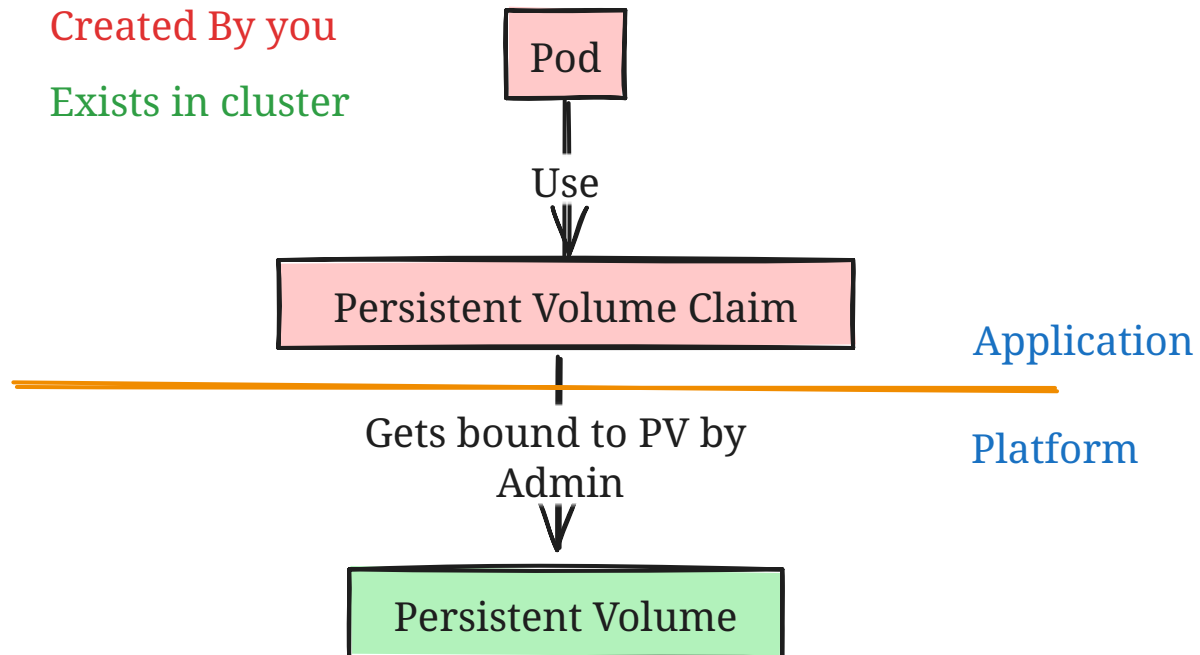
- Settings for dynamically provisioning a PV from a PVC

DYNAMIC PROVISIONING



STATIC PROVISIONING

Created By you
Exists in cluster



ATTACHMENT SEMANTICS

- ReadWriteOnce / RWO - Attached only at one Node at a time
- ReadWriteMany / RWX - Attached to multiple Nodes at a time
- ReadOnlyMany / ROX - Attached to multiple Nodes at a time (ReadOnly)
- ReadWriteOncePod / RWOP - Attached only at one **Pod** at a time

TYPES OF STORAGE PROVIDERS

- Cloud Provider Integrations
 - AWS EBS+EFS, Azure Files
- Connection to external storage solutions
 - Pure Storage, SMB, VSphere, JuiceFS
- Storage Operators
 - Rook Ceph, Longhorn
- Special use case
 - Local Path, Secrets CSI Driver

List of known ones: <https://kubernetes-csi.github.io/docs/drivers.html>

WHAT IS CSI?

CSI DRIVERS?

- Maps storage solution to Kubernetes
- Handles volume lifecycle
 - Provisioning, Attachment, Mounting, Deleting
- Maintained by storage solution provider*
- Run as Pods in Kubernetes*
- Installed via Helm Charts/K8s manifests or comes with your cluster

VERY SHORT HISTORY OF KUBERNETES STORAGE

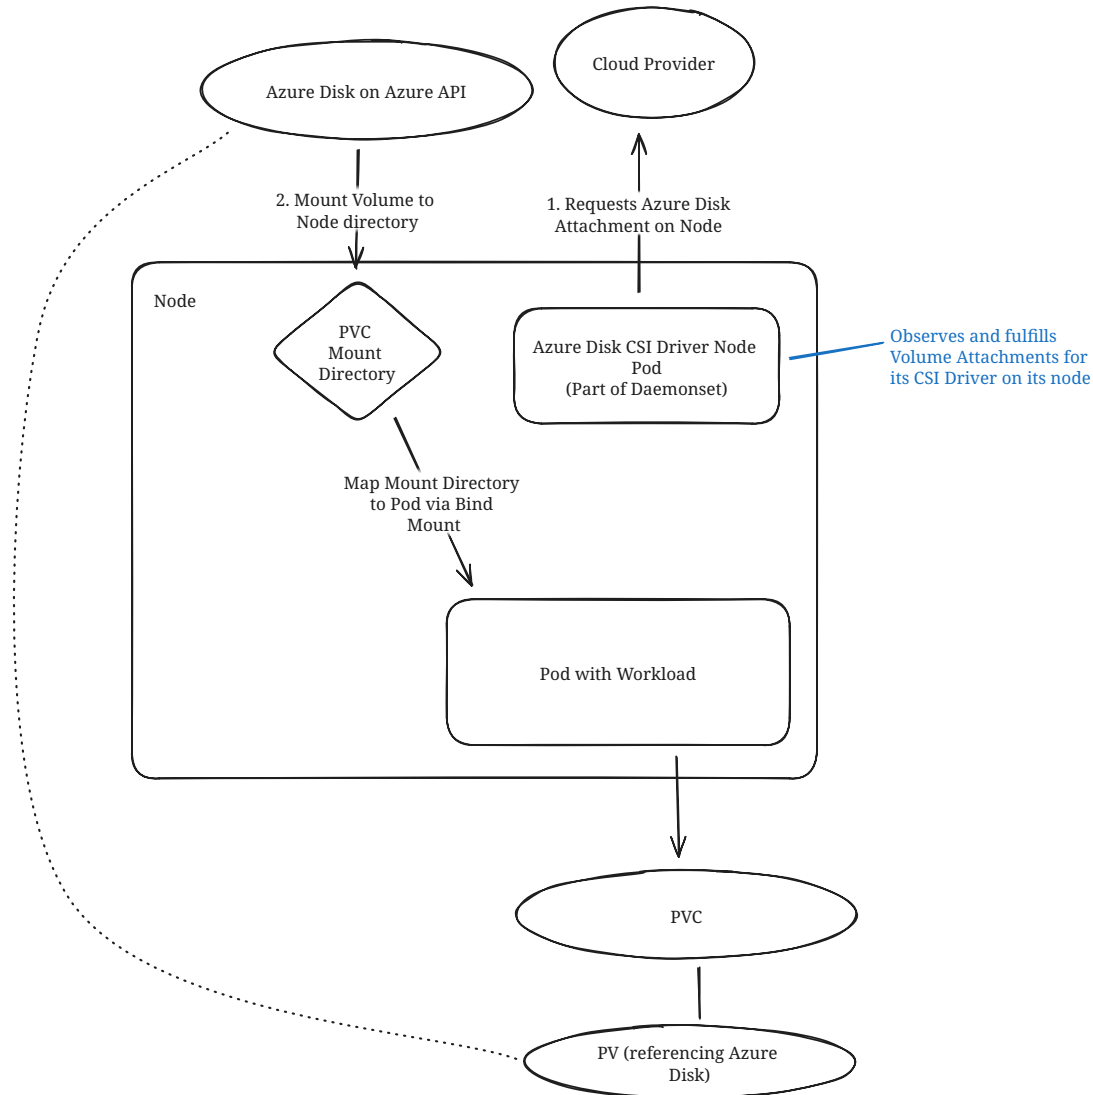
- In the beginning: in-tree Storage Providers
 - Since 1.2: FlexVolume for out-of-tree providers
 - Since 1.9: Container Storage Interface introduced
 - 1.13: CSI is GA
 - Since around 1.21: in-tree Storage Providers deprecation
 - Since 1.23: FlexVolume deprecated
- > CSI is K8s Storage today*

OPTIONAL FEATURES

- Online Volume Expansion
- Volume Cloning/Snapshotting
- Raw Block Volumes

HOW CSI MOUNTING WORKS?

EXAMPLE: ATTACHING AN AZURE DISK VOLUME



**HOW TO SELECT THE
RIGHT PROVIDER FOR
YOU?**

Everybody gangsta until the storage cluster stops
clustering

-They

- the easiest problems are the ones you do not solve yourself
- Explore available options
- Use a cloud provider integration, if you can
- Most SANs have a CSI Provider (PureStorage, NetApp, etc.)
- Use well established solution with a high degree of automation, like Rook Ceph if you have to do it yourself

BACKUPS

Snapshots are not Backups

Replicas are not Backups

One Backup is no Backup

Untested Backups do not exist

If you can delete it using the credentials used to create
it, it does not exist

A GOOD BACKUP IS

- Immutable
- Tested
- Replicated
- Infrastructure independent

=> Object Storage

HOW TO SETUP?

- Create Bucket outside of regular infrastructure permission system
- Create Credentials without delete permissions, only hiding for deployment into infrastructure
- Use a Lifecycle Rule to delete backups after they have been hidden for a while (90 days - x Years)

VELERO

- Backups K8s resources as well
 - Please use GitOps for that
- Difficult to integrate with anything else than plain PVCs*
- Can do a CSI Snapshot and extract data to an object store
- Does its operations using a DaemonSet

K8UP

- Back up all PVCs into an Object Store using Restic
- Manages helper pods for Application Level Backups
- No cluster level configuration necessary
- Does backups using Jobs -> only allocating resources during operations
- One Backup Custom Resource per namespace

```
apiVersion: k8up.io/v1
kind: Schedule
metadata:
  name: confluence
spec:
  backup:
    schedule: "@daily-random"
    podSecurityContext:
      runAsUser: 0
  check:
    schedule: "@weekly-random"
  prune:
    schedule: "@weekly-random"
    retention:
      keepDaily: 30
      keepMonthly: 6
      keepWeekly: 8
  backend:
    s3:
      endpoint: https://s3.eu-central-003.backblazeb2.com
      bucket: "backupbucket/apps-confluence"
      accessKeyIDSecretRef:
        name: k8up-credentials
        key: username
      secretAccessKeySecretRef:
```

```
apiVersion: k8up.io/v1
kind: PreBackupPod
metadata:
  name: confluence-db-backup
spec:
  backupCommand: pg_dump
  fileExtension: .sql
  pod:
    spec:
      containers:
      - name: confluence-db-dump
        image: postgres:latest
        env:
          - name: PGDATABASE
            value: confluence
          - name: PGHOST
            value: acid-confluence
          - name: PGPASSWORD
            valueFrom:
              secretKeyRef:
                key: password
                name: confluence.acid-confluence.credentials.postgresql.acid
          - name: PGUSER
            valueFrom:
              secretKeyRef:
```

ALTERNATIVES TO PVCS

S3 STYLE OBJECT STORES

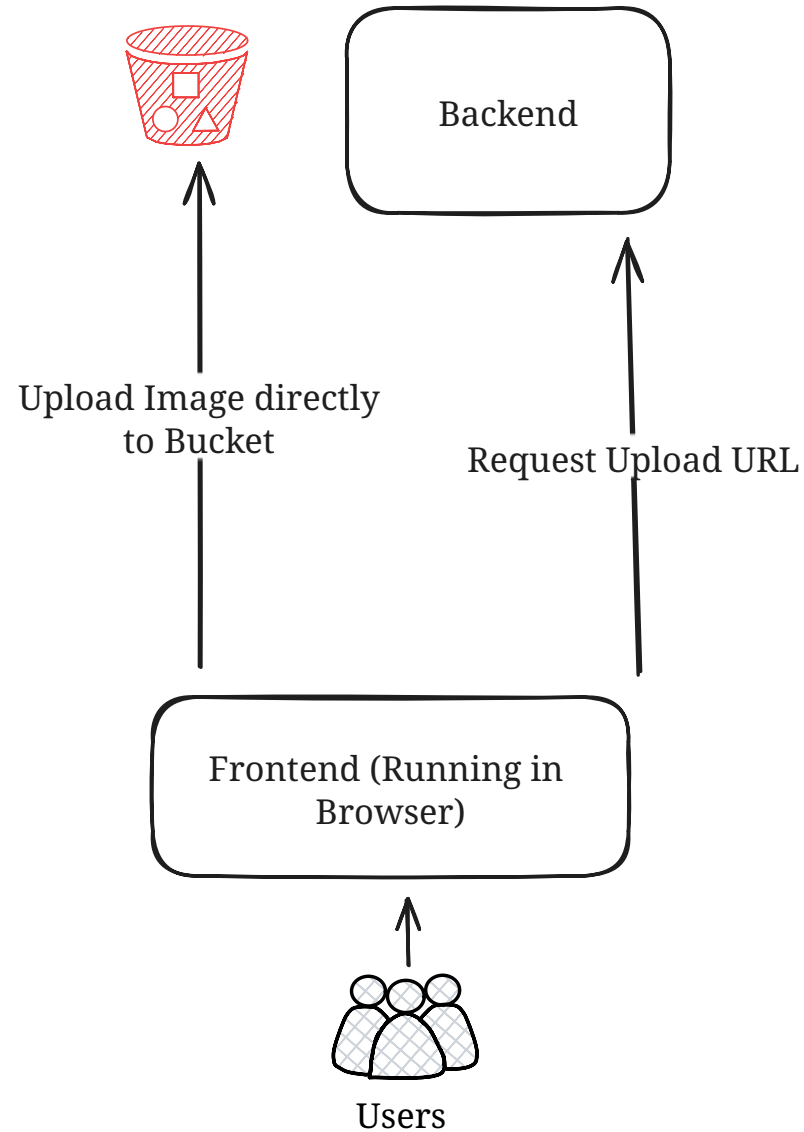
- Read-After-Write consistent*
- High Time-To-First-Byte Latency (100ms+)*
- Partial Downloads
- Per-Request Performance limit instead of Per-Bucket limit
- Big Limitation: NO PARTIAL OBJECT UPDATES
- Cheaper than block storage*

PRESIGNED URLs IN S3

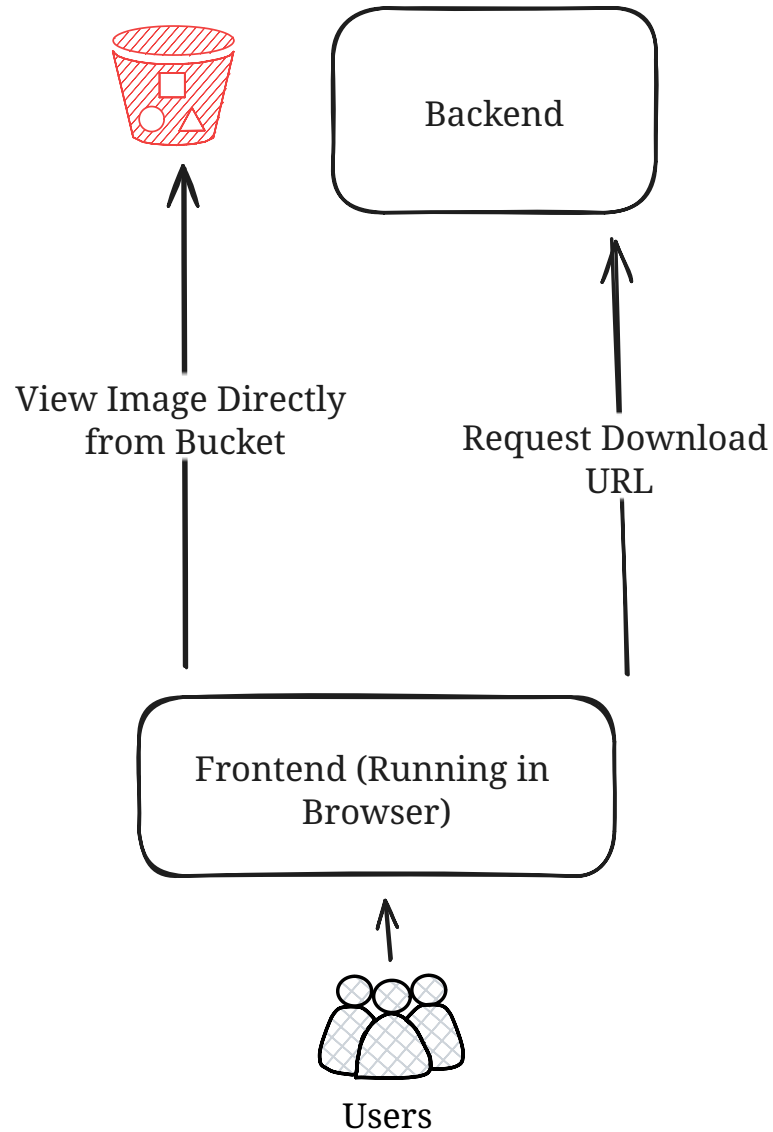
Delegate temporary S3 Permissions directly to the end-user client

Allows you to move data without passing through your own application/infrastructure

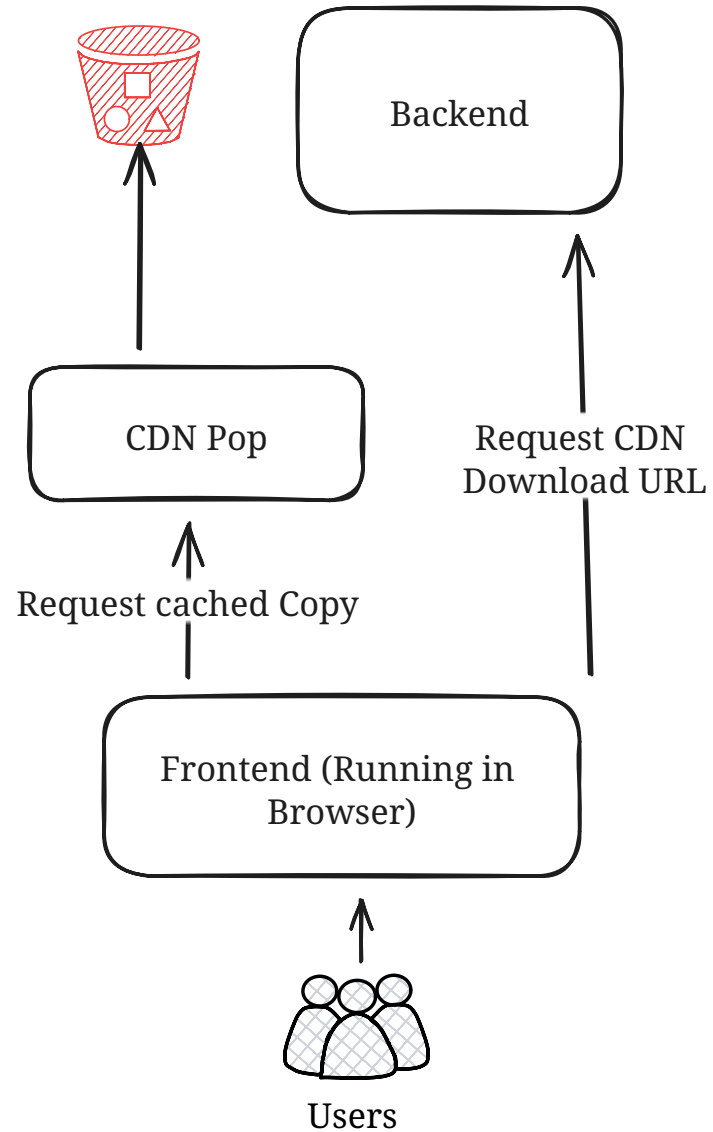
UPLOAD



DOWNLOAD



USING A CDN



RUNNING DATABASES IN-CLUSTER

WHY?

WHY?

Managed Databases are fine, but

- Poor integration with Kubernetes CD
- Expensive
- Different features across cloud vendors
- Latency
- Separate Database Teams add communication overhead

GENERAL ADVICE (REGARDLESS OF DB)

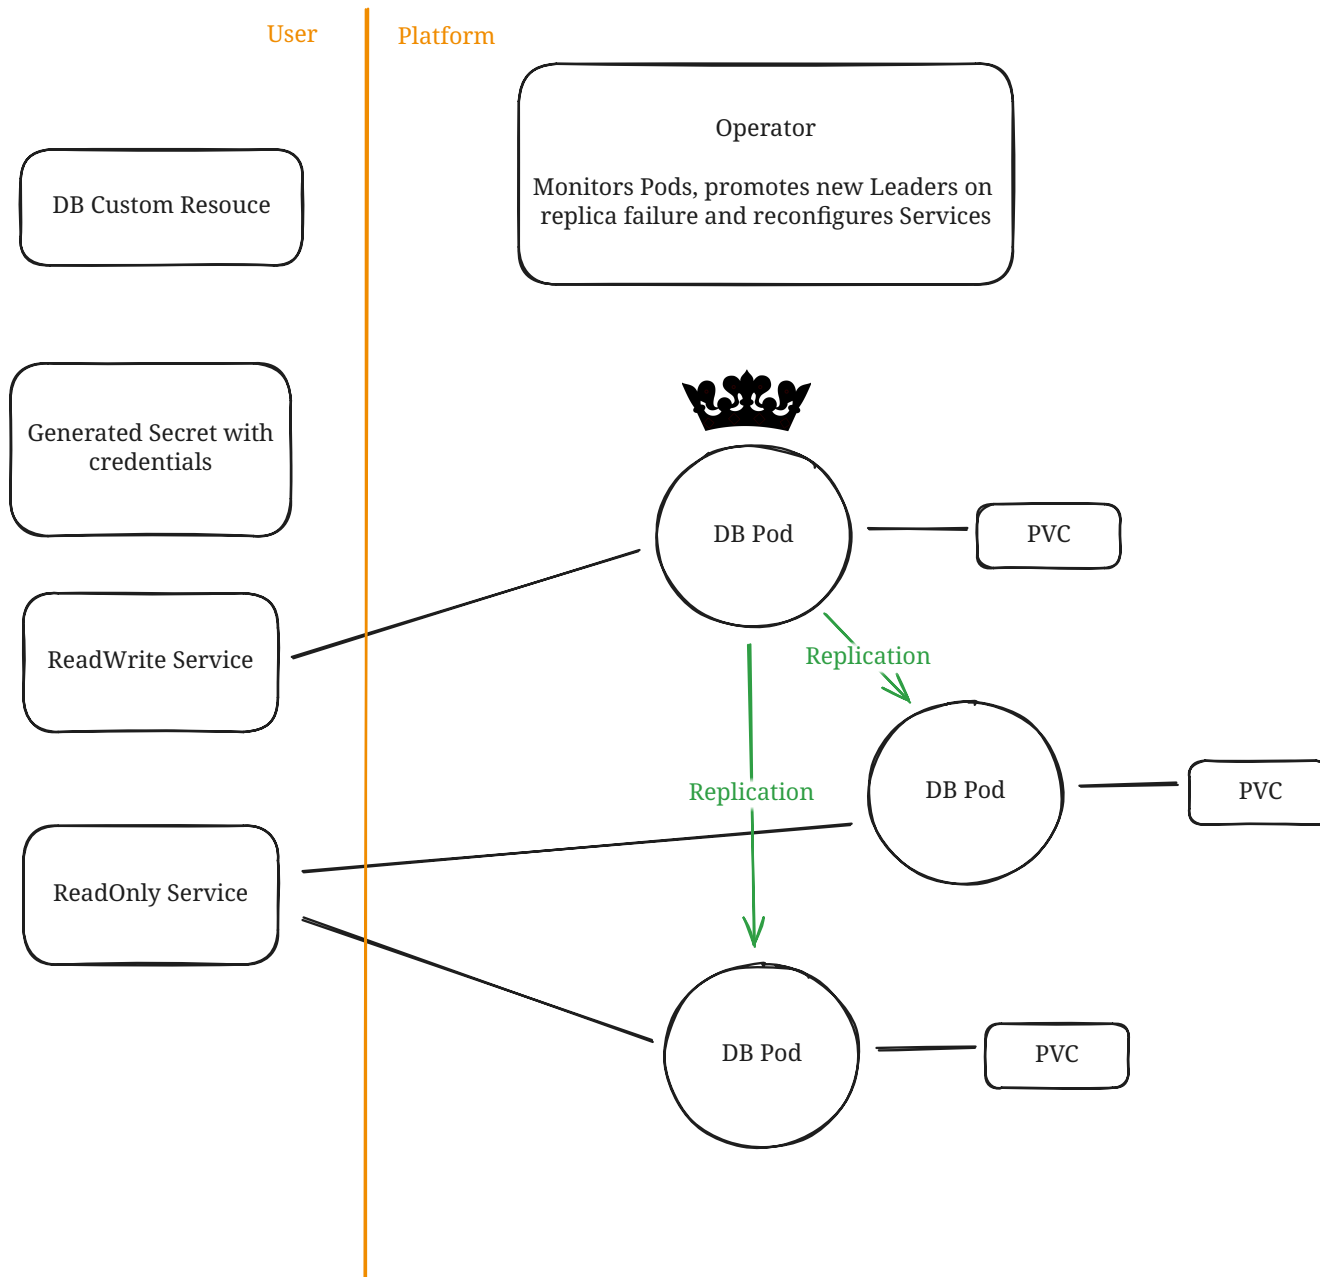
- Run a DB cluster, please
- Use an operator if you need persistence
- Test your application to handle DB reconnection
- Use monitoring and backups

ACID DATABASES

Postgres, MySQL, MariaDB

ACID DATABASES CONSIDERATIONS

- Use a fast, low latency PVC
- Run your DB in the same K8s cluster as the application



CLOUDNATIVE PG

CNCF Project

Does all the things

WAL Streaming, Automatic Backups, Replication,
Upgrades, Point-In-Time Restore

<https://cloudnative-pg.io/>

ZALANDO POSTGRES OPERATOR

A bit older, based on a project called Patroni for Docker

Similar to CNPG, but some features hard to use outside of AWS (e.g. WAL Streaming)

Use the more modern CNPG

<https://github.com/zalando/postgres-operator>

MYSQL OPERATOR

Built by Oracle

Tried to use it, but had issues all around in the
beginning of 2023

Might be better now

¬_(ツ)_/¬

VITESS

- MySQL Compatible
- Kubernetes Native design
- Supports data sharding
- Some compatibility limitations: <https://vitess.io/docs/19.0/reference/compatibility/mysql-compatibility/>

ORACLEDB OPERATOR

<https://github.com/oracle/oracle-database-operator>

BACKUPS

- WAL Archiving
- Regular Dumps, if size allows for it

**FUN TIMES WITH
NETWORK ATTACHED
STORAGE**

RUNNING STORAGE CLUSTERS INSIDE KUBERNETES

Long: Failover on the protocol level is a must

Short: Yes, if it is Rook Ceph

LATENCY

- Using networked filesystems introduces significant IO latency >0.5ms
- Limits your sequential throughput to 2000 ops/s
- Leads to high commit/flush latency

Solution:

Agressively parallelize your IO - Your throughput will scale pretty well*

THANK YOU

